

Heart Disease Prediction using machine learning

Table of Contents

1. Import Packages
2. Data Preprocessing
3. Encoding Categorical Data
4. EDA
5. Preparing ML models
6. Models evaluation
7. Conclusion

Packages Required

```
In [1]: #Loading dataset
import pandas as pd
import numpy as np
#visualisation
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
#EDA
from collections import Counter
import ydata_profiling as pp
# data preprocessing
from sklearn.preprocessing import StandardScaler, LabelEncoder
# data splitting
from sklearn.model_selection import train_test_split
# data modeling
from sklearn.metrics import confusion_matrix, accuracy_score, roc_curve, classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from xgboost import XGBClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.decomposition import PCA
from sklearn.svm import SVC
```

```
#ensembling
from mlxtend.classifier import StackingCVClassifier
```

```
In [2]: data = pd.read_csv('heart.csv')
data.head()
```

```
Out[2]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	ti
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	

```
In [3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1025 non-null   int64
1   sex         1025 non-null   int64
2   cp          1025 non-null   int64
3   trestbps    1025 non-null   int64
4   chol        1025 non-null   int64
5   fbs         1025 non-null   int64
6   restecg     1025 non-null   int64
7   thalach     1025 non-null   int64
8   exang       1025 non-null   int64
9   oldpeak     1025 non-null   float64
10  slope       1025 non-null   int64
11  ca          1025 non-null   int64
12  thal        1025 non-null   int64
13  target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

Data Preprocessing

```
In [4]: # Checking for missing values
print(data.isnull().sum())

# Dropping duplicate rows
data = data.drop_duplicates()
```

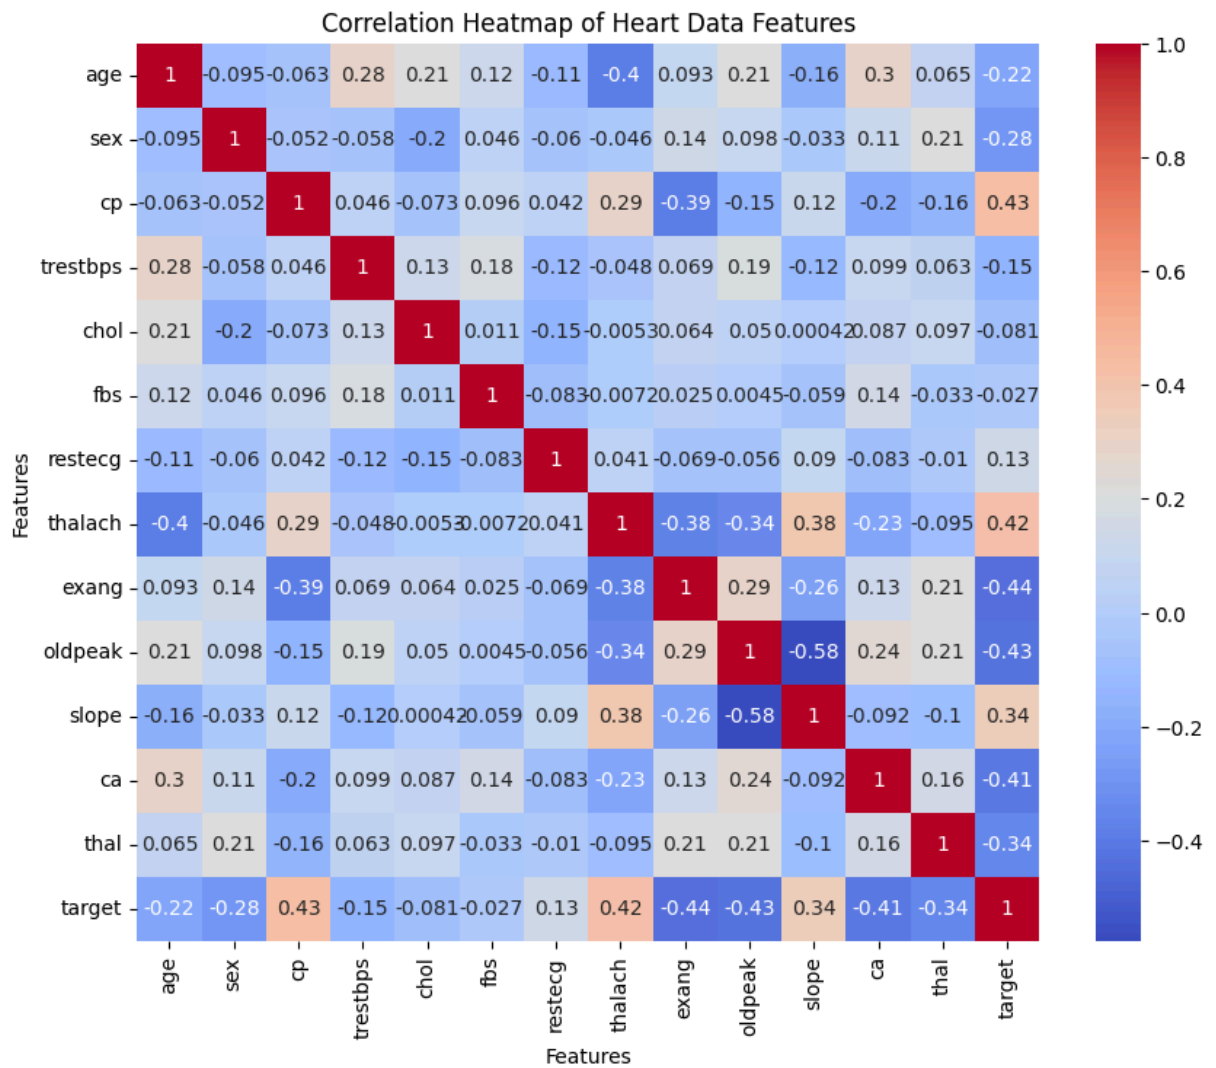
```
age          0
sex          0
cp          0
trestbps    0
chol        0
fbs         0
restecg     0
thalach     0
exang       0
oldpeak     0
slope       0
ca          0
thal        0
target      0
dtype: int64
```

Encoding Categorical Data

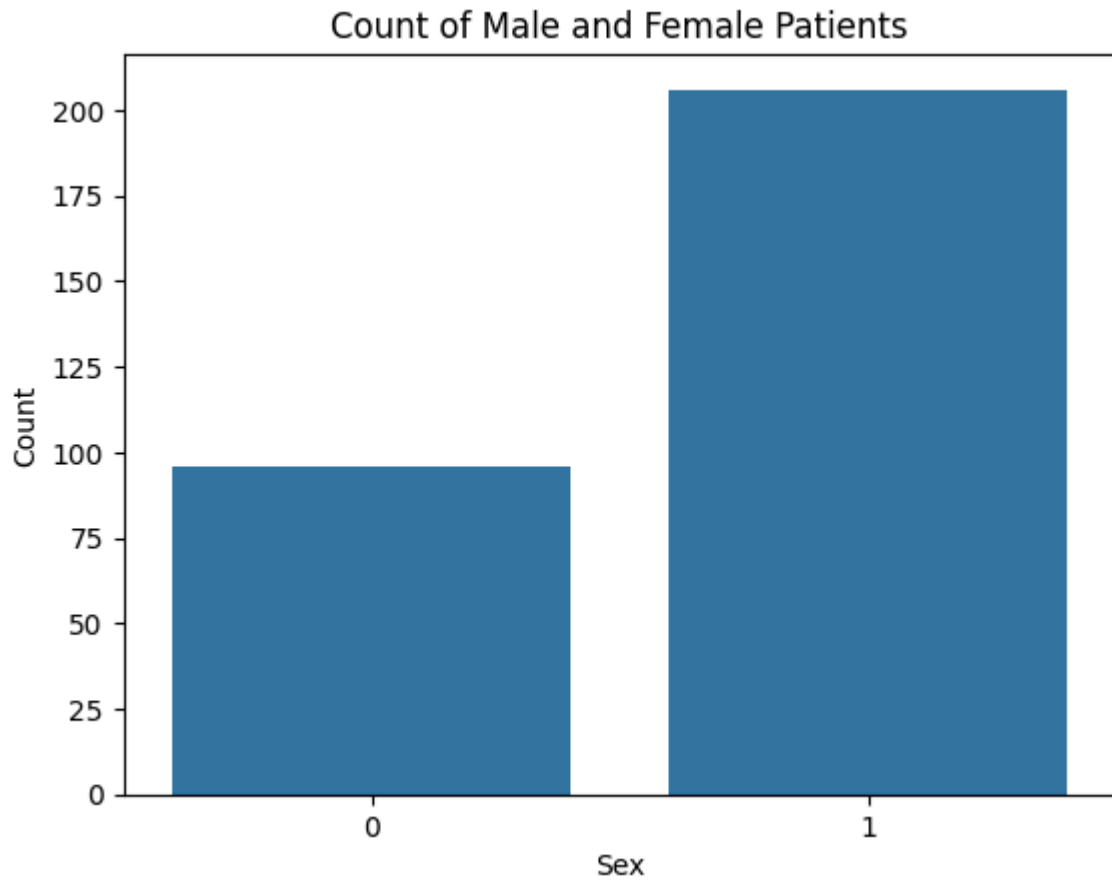
```
In [5]: # Encoding categorical features
categorical_columns = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal']
for col in categorical_columns:
    le = LabelEncoder()
    data[col] = le.fit_transform(data[col])
```

EDA

```
In [6]: # Correlation heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(data.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap of Heart Data Features')
plt.xlabel('Features')
plt.ylabel('Features')
plt.show()
```




```
In [7]: # Countplot for the 'sex' feature (without palette)
sns.countplot(x='sex', data=data)
plt.title('Count of Male and Female Patients')
plt.xlabel('Sex')
plt.ylabel('Count')
plt.show()
```




```
In [8]: # Histogram for feature distribution
numerical_cols = ['age', 'trestbps', 'chol', 'thalach', 'oldpeak']
for col in numerical_cols:
    plt.figure(figsize=(8, 4))
    sns.histplot(data[col], kde=True)
    plt.title(f'Distribution of {col}')
    plt.xlabel(f'{col}')
    plt.ylabel('Frequency')
    plt.show()
```


Distribution of age



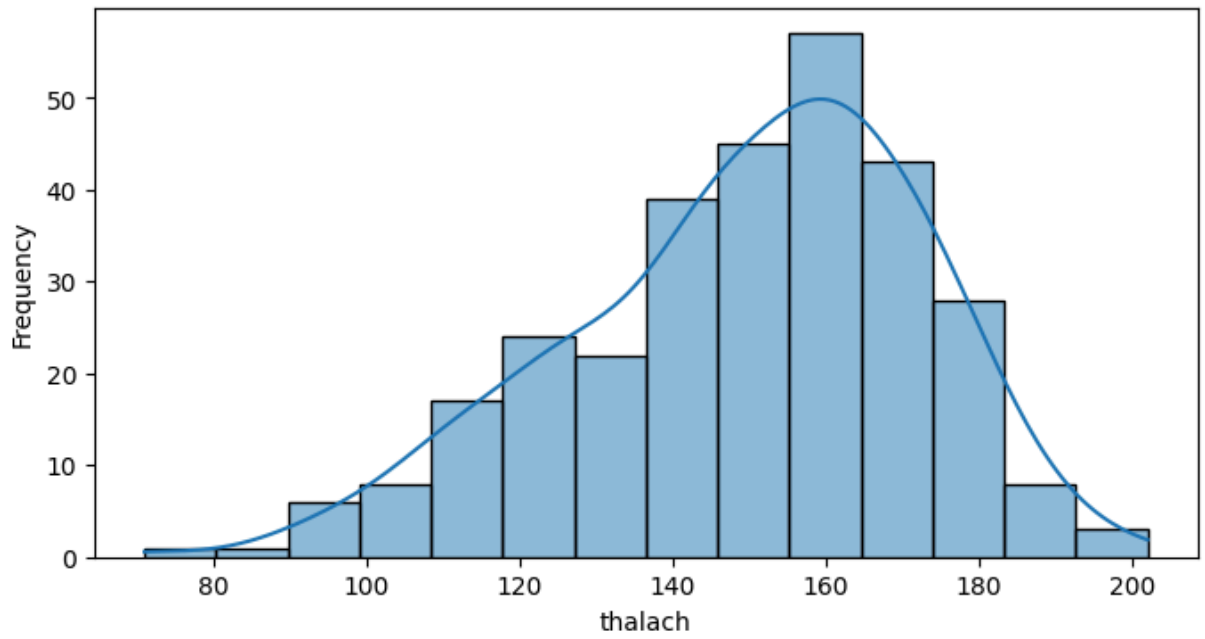
Distribution of trestbps

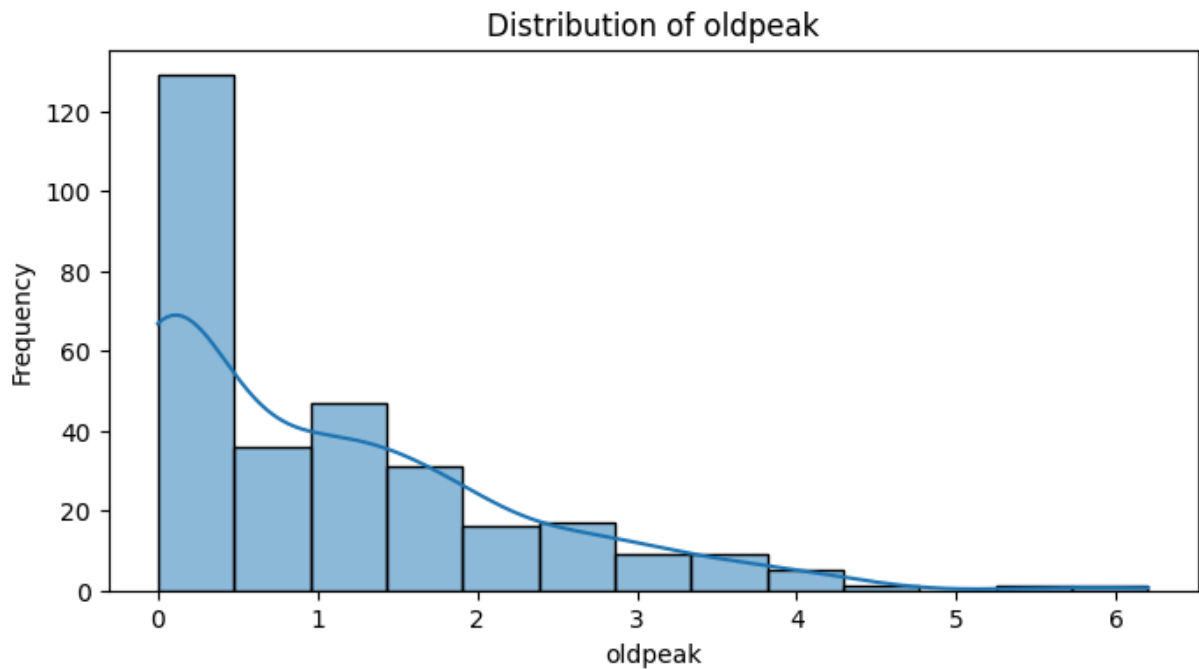


Distribution of chol



Distribution of thalach





Model prepration

```
In [9]: y = data["target"]
X = data.drop('target',axis=1)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_st
```

Before applying algorithm we should check whether the data is equally splitted or not, because if data is not splitted equally it will cause for data imbalacing problem

```
In [10]: print(y_test.unique())
Counter(y_train)
```

```
[1 0]
```

```
Out[10]: Counter({1: 128, 0: 113})
```

```
In [11]: scaler = StandardScaler()
scaled_data = scaler.fit_transform(data.drop('target', axis=1))
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```


Principal Component Analysis (PCA)

```
In [12]: # Applying PCA and selecting PC1 and PC2
pca = PCA(n_components=2)
X_pca = pca.fit_transform(scaled_data)

# Visualizing PCA results
plt.figure(figsize=(8, 6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y, cmap='coolwarm')
plt.xlabel('PC1')
```



```
plt.ylabel('PC2')
plt.title('PCA: PC1 vs PC2')
plt.show()
```



ML models

Here I take different machine learning algorithm and try to find algorithm which predict accurately.

1. Logistic Regression
2. Naive Bayes
3. Random Forest Classifier
4. Extreme Gradient Boost
5. K-Nearest Neighbour
6. Decision Tree
7. Support Vector Machine

```
In [13]: m1 = 'Logistic Regression'
lr = LogisticRegression()
model = lr.fit(X_train, y_train)
lr_predict = lr.predict(X_test)
lr_conf_matrix = confusion_matrix(y_test, lr_predict)
lr_acc_score = accuracy_score(y_test, lr_predict)
```

```

print("confussion matrix")
print(lr_conf_matrix)
print("\n")
print("Accuracy of Logistic Regression:",lr_acc_score*100,'\n')
print(classification_report(y_test,lr_predict))

```

```

confussion matrix
[[20  5]
 [ 5 31]]

```

Accuracy of Logistic Regression: 83.60655737704919

	precision	recall	f1-score	support
0	0.80	0.80	0.80	25
1	0.86	0.86	0.86	36
accuracy			0.84	61
macro avg	0.83	0.83	0.83	61
weighted avg	0.84	0.84	0.84	61

```

In [14]: m2 = 'Naive Bayes'
nb = GaussianNB()
nb.fit(X_train,y_train)
nbpred = nb.predict(X_test)
nb_conf_matrix = confusion_matrix(y_test, nbpred)
nb_acc_score = accuracy_score(y_test, nbpred)
print("confussion matrix")
print(nb_conf_matrix)
print("\n")
print("Accuracy of Naive Bayes model:",nb_acc_score*100,'\n')
print(classification_report(y_test,nbpred))

```

```

confussion matrix
[[20  5]
 [ 7 29]]

```

Accuracy of Naive Bayes model: 80.32786885245902

	precision	recall	f1-score	support
0	0.74	0.80	0.77	25
1	0.85	0.81	0.83	36
accuracy			0.80	61
macro avg	0.80	0.80	0.80	61
weighted avg	0.81	0.80	0.80	61

```

In [15]: m3 = 'Random Forest Classifier'
rf = RandomForestClassifier(n_estimators=20, random_state=12,max_depth=5)
rf.fit(X_train,y_train)
rf_predicted = rf.predict(X_test)
rf_conf_matrix = confusion_matrix(y_test, rf_predicted)

```

```

rf_acc_score = accuracy_score(y_test, rf_predicted)
print("confussion matrix")
print(rf_conf_matrix)
print("\n")
print("Accuracy of Random Forest:", rf_acc_score*100, '\n')
print(classification_report(y_test, rf_predicted))

```

confussion matrix

```

[[18  7]
 [ 5 31]]

```

Accuracy of Random Forest: 80.32786885245902

	precision	recall	f1-score	support
0	0.78	0.72	0.75	25
1	0.82	0.86	0.84	36
accuracy			0.80	61
macro avg	0.80	0.79	0.79	61
weighted avg	0.80	0.80	0.80	61

```

In [16]: m4 = 'Extreme Gradient Boost'
xgb = XGBClassifier(learning_rate=0.01, n_estimators=25, max_depth=15, gamma=0.6, su
               reg_lambda=2, booster='dart', colsample_bylevel=0.6, colsample_
xgb.fit(X_train, y_train)
xgb_predicted = xgb.predict(X_test)
xgb_conf_matrix = confusion_matrix(y_test, xgb_predicted)
xgb_acc_score = accuracy_score(y_test, xgb_predicted)
print("confussion matrix")
print(xgb_conf_matrix)
print("\n")
print("Accuracy of Extreme Gradient Boost:", xgb_acc_score*100, '\n')
print(classification_report(y_test, xgb_predicted))

```

confussion matrix

```

[[13 12]
 [ 0 36]]

```

Accuracy of Extreme Gradient Boost: 80.32786885245902

	precision	recall	f1-score	support
0	1.00	0.52	0.68	25
1	0.75	1.00	0.86	36
accuracy			0.80	61
macro avg	0.88	0.76	0.77	61
weighted avg	0.85	0.80	0.79	61

```

In [17]: m5 = 'K-NeighborsClassifier'
knn = KNeighborsClassifier(n_neighbors=10)
knn.fit(X_train, y_train)

```

```

knn_predicted = knn.predict(X_test)
knn_conf_matrix = confusion_matrix(y_test, knn_predicted)
knn_acc_score = accuracy_score(y_test, knn_predicted)
print("confussion matrix")
print(knn_conf_matrix)
print("\n")
print("Accuracy of K-NeighborsClassifier:",knn_acc_score*100,'\n')
print(classification_report(y_test,knn_predicted))

```

```

confussion matrix
[[19  6]
 [ 5 31]]

```

Accuracy of K-NeighborsClassifier: 81.9672131147541

	precision	recall	f1-score	support
0	0.79	0.76	0.78	25
1	0.84	0.86	0.85	36
accuracy			0.82	61
macro avg	0.81	0.81	0.81	61
weighted avg	0.82	0.82	0.82	61

```

In [18]: m6 = 'DecisionTreeClassifier'
dt = DecisionTreeClassifier(criterion = 'entropy',random_state=0,max_depth = 6)
dt.fit(X_train, y_train)
dt_predicted = dt.predict(X_test)
dt_conf_matrix = confusion_matrix(y_test, dt_predicted)
dt_acc_score = accuracy_score(y_test, dt_predicted)
print("confussion matrix")
print(dt_conf_matrix)
print("\n")
print("Accuracy of DecisionTreeClassifier:",dt_acc_score*100,'\n')
print(classification_report(y_test,dt_predicted))

```

```

confussion matrix
[[18  7]
 [ 6 30]]

```

Accuracy of DecisionTreeClassifier: 78.68852459016394

	precision	recall	f1-score	support
0	0.75	0.72	0.73	25
1	0.81	0.83	0.82	36
accuracy			0.79	61
macro avg	0.78	0.78	0.78	61
weighted avg	0.79	0.79	0.79	61

```

In [19]: m7 = 'Support Vector Classifier'
svc = SVC(kernel='rbf', C=2)

```

```

svc.fit(X_train, y_train)
svc_predicted = svc.predict(X_test)
svc_conf_matrix = confusion_matrix(y_test, svc_predicted)
svc_acc_score = accuracy_score(y_test, svc_predicted)
print("confussion matrix")
print(svc_conf_matrix)
print("\n")
print("Accuracy of Support Vector Classifier:",svc_acc_score*100,'\n')
print(classification_report(y_test,svc_predicted))

```

confussion matrix

```

[[18  7]
 [ 6 30]]

```


Accuracy of Support Vector Classifier: 78.68852459016394

	precision	recall	f1-score	support
0	0.75	0.72	0.73	25
1	0.81	0.83	0.82	36
accuracy			0.79	61
macro avg	0.78	0.78	0.78	61
weighted avg	0.79	0.79	0.79	61

```

In [20]: imp_feature = pd.DataFrame({'Feature': ['age', 'sex', 'cp', 'trestbps', 'chol', 'fb
        'exang', 'oldpeak', 'slope', 'ca', 'thal'], 'Importance': xgb.feature_import
plt.figure(figsize=(10,4))
plt.title("barplot Represent feature importance ")
plt.xlabel("importance ")
plt.ylabel("features")
plt.barh(imp_feature['Feature'], imp_feature['Importance'], color=['red', 'green',
plt.show()

```



```

In [21]: lr_false_positive_rate,lr_true_positive_rate,lr_threshold = roc_curve(y_test,lr_pre
nb_false_positive_rate,nb_true_positive_rate,nb_threshold = roc_curve(y_test,nbpred
rf_false_positive_rate,rf_true_positive_rate,rf_threshold = roc_curve(y_test,rf_pre
xgb_false_positive_rate,xgb_true_positive_rate,xgb_threshold = roc_curve(y_test,xgb


```

```

knn_false_positive_rate,knn_true_positive_rate,knn_threshold = roc_curve(y_test,knn
dt_false_positive_rate,dt_true_positive_rate,dt_threshold = roc_curve(y_test,dt_pre
svc_false_positive_rate,svc_true_positive_rate,svc_threshold = roc_curve(y_test,svc

sns.set_style('whitegrid')
plt.figure(figsize=(10,5))
plt.title('Reciver Operating Characterstic Curve')
plt.plot(lr_false_positive_rate,lr_true_positive_rate,label='Logistic Regression')
plt.plot(nb_false_positive_rate,nb_true_positive_rate,label='Naive Bayes')
plt.plot(rf_false_positive_rate,rf_true_positive_rate,label='Random Forest')
plt.plot(xgb_false_positive_rate,xgb_true_positive_rate,label='Extreme Gradient Boo
plt.plot(knn_false_positive_rate,knn_true_positive_rate,label='K-Nearest Neighbor')
plt.plot(dt_false_positive_rate,dt_true_positive_rate,label='Desion Tree')
plt.plot(svc_false_positive_rate,svc_true_positive_rate,label='Support Vector Class
plt.plot([0,1],ls='--')
plt.plot([0,0],[1,0],c='.5')
plt.plot([1,1],c='.5')
plt.ylabel('True positive rate')
plt.xlabel('False positive rate')
plt.legend()
plt.show()

```



Model Evaluation

```

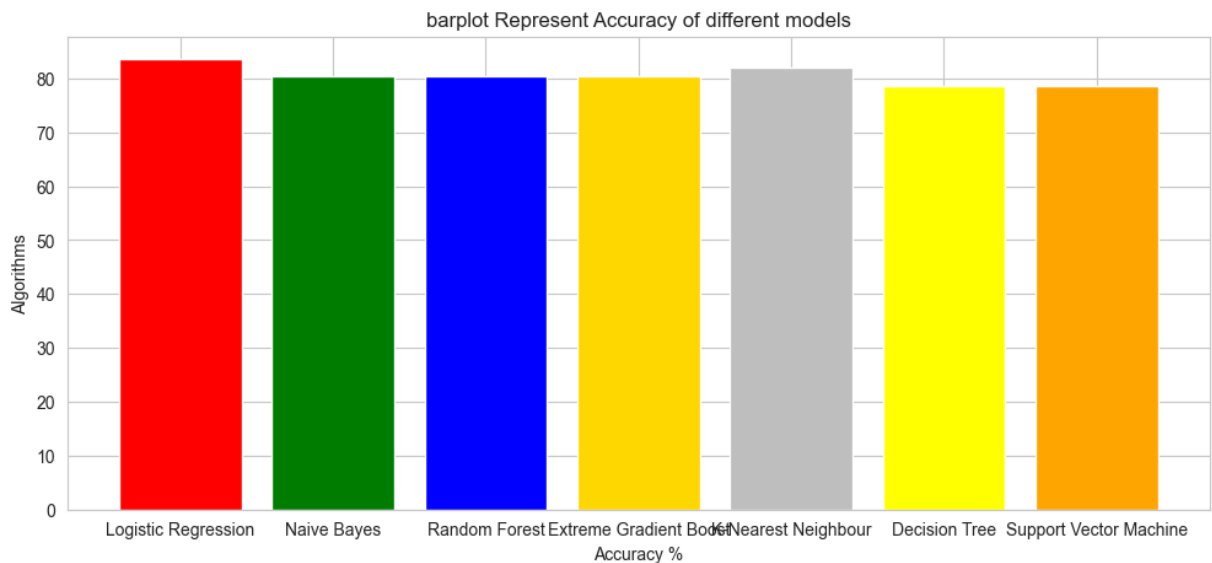
In [22]: model_ev = pd.DataFrame({'Model': ['Logistic Regression','Naive Bayes','Random Fore
        'K-Nearest Neighbour','Decision Tree','Support Vector Machine']
        nb_acc_score*100,rf_acc_score*100,xgb_acc_score*100,knn_acc_sco
model_ev

```

Out[22]:

	Model	Accuracy
0	Logistic Regression	83.606557
1	Naive Bayes	80.327869
2	Random Forest	80.327869
3	Extreme Gradient Boost	80.327869
4	K-Nearest Neighbour	81.967213
5	Decision Tree	78.688525
6	Support Vector Machine	78.688525

```
In [23]: colors = ['red', 'green', 'blue', 'gold', 'silver', 'yellow', 'orange',]
plt.figure(figsize=(12,5))
plt.title("barplot Represent Accuracy of different models")
plt.xlabel("Accuracy %")
plt.ylabel("Algorithms")
plt.bar(model_ev['Model'],model_ev['Accuracy'],color = colors)
plt.show()
```



Conclusion

In this project, we applied various machine learning models to predict heart disease based on a dataset containing patient information. After performing data preprocessing, including scaling and encoding categorical variables, we utilized **Principal Component Analysis (PCA)** to reduce the dimensionality and visualize the data using the first two principal components (PC1 and PC2).

Several machine learning models were trained and evaluated, including:

- Logistic Regression

- Naive Bayes
- Random Forest Classifier
- Extreme Gradient Boosting (XGBoost)
- K-Nearest Neighbors (KNN)
- Decision Tree Classifier
- Support Vector Machine (SVM)

Model Performance:

The accuracy results of the models were as follows:

1. **Logistic Regression:** 83.61%
2. **Naive Bayes:** 80.33%
3. **Random Forest Classifier:** 80.33%
4. **Extreme Gradient Boosting:** 80.33%
5. **K-Nearest Neighbors (KNN):** 81.97%
6. **Decision Tree Classifier:** 78.69%
7. **Support Vector Machine (SVM):** 78.69%

Insights:

- Logistic Regression emerged as the most accurate model, closely followed by K-Nearest Neighbors.
- All models performed similarly, with accuracies between 78% and 83%, indicating consistent predictive power across different techniques.
- The models demonstrated that features like chest pain type (cp), maximum heart rate achieved (thalach), and ST depression (oldpeak) are important predictors for heart disease.
- PCA helped to visualize the separation between classes, confirming that the data contains patterns that machine learning models can exploit effectively.